

BEANSTECH

Trustworthy AI for Regulated Sectors

Legal · Health · Finance — sovereign, citation-grounded intelligence

NVIDIA Inception Program Application — 2026

THE PROBLEM

In regulated sectors, wrong AI output is a legal liability

Hallucination

Generic LLMs invent case law, diagnoses and rules — unusable where citations are mandatory.

Data sovereignty

Client data cannot flow through foreign clouds: LGPD, CFM 2.454, BACEN 3.978 constrain processing.

Unreadable sources

Brazilian court systems produce millions of PDFs with corrupted fonts that no AI can read.

Result: law firms, hospitals and financial institutions in Brazil cannot adopt mainstream AI — **the sectors that need trustworthy AI the most are the ones served the least.**

THE SOLUTION

One sovereign AI platform, three regulated verticals

Legal AI

RAG over 65M court decisions with citation-grounded answers and INPI-registered anti-hallucination protocol. Platforms: RAGJur, Advogando.AI, Minuta.Tech.

Health AI

7 CFM-compliant platforms (FHIR R4): clinical AI scribe, exam support, digital prescription — compliance for the Aug/2026 CFM deadline.

RegTech

Production AML engine — 25 detection typologies (BACEN 3.978) serving banks and public institutions.

Every answer cites its source. Every byte stays in Brazil. Zero telemetry by design.

TRACTION

Production systems, institutional validation

65M+

court decisions indexed

55

tribunals covered daily

15+

live platforms

160

specialized AI agents

- ▶ **2 large clients in production** — legal and financial institutions.
- ▶ **Hospital Albert Einstein** — innovation team formally evaluating our architecture (their initiative).
- ▶ **Google for Startups 2025 · Anthropic Partner** — ecosystem recognitions; Oracle Startup program application submitted.
- ▶ **Public-sector offer** — technology-transfer proposal to Brazil's Federal Prosecution Office under formal review.

DIFFERENTIATORS

Why institutions trust the stack

- ▶ **Citation-verified outputs** — INPI-registered anti-hallucination protocol; every legal claim is checked against the source decision (0 false citations tolerated).
- ▶ **Data sovereignty** — on-premise or Brazil-region inference, zero telemetry: no token, metric or document ever leaves the institution's network.
- ▶ **Compliance by design** — CFM 2.454 (health), BACEN 3.978 (AML), LGPD mapped per vertical before feature one.
- ▶ **Ingestion that reads Brazil** — OCR engine that decodes corrupted court PDFs (PJe/e-SAJ); a 158-page case file is digitized in seconds.
- ▶ **Founder-market fit** — lawyer (OAB) + Google Cloud Digital Leader building for the sectors he practiced in.

Where CUDA runs our stack today — and tomorrow

TODAY (production)

- ▶ Open-weight LLM inference (Qwen/GLM, quantized) on NVIDIA L4/A100 (GCP, Brazil region)
- ▶ CUDA-accelerated embeddings + reranking over tens of millions of judgments
- ▶ Multimodal pipelines (document/image) on GPU workers

ROADMAP WITH INCEPTION

- ▶ TensorRT-LLM + vLLM: 2-4x latency/cost gains for legal drafting workloads
- ▶ Triton Inference Server for multi-tenant institutional serving
- ▶ On-prem GPU appliances (zero-telemetry deployments at courts, hospitals, banks)
- ▶ Open-weight video generation (Wan) for institutional training content

Brazil's regulated sectors, underserved by global AI

- ▶ **Legal** — 1.3M+ lawyers, 100M+ new case pages/year; citation duty makes generic LLMs legally unusable.
- ▶ **Health** — CFM resolution 2.454 deadline (Aug/2026) forces every physician onto compliant tooling.
- ▶ **Finance** — BACEN Circular 3.978 obliges AML typology coverage; mid-size institutions lack in-house AI.
- ▶ **wedge** — we index what others can't read (court PDFs) and verify what others invent (citations).

Regulated-domain depth + production engineering

Matheus Feijão — Founder & CEO

- ▶ Lawyer (OAB/DF) — 13 years in federal justice system
- ▶ Google Cloud Digital Leader · Google for Startups founder
- ▶ Anthropic Partner · GLM/Qwen production deployments
- ▶ M.Sc. candidate (Digital Law, IDP-SP)

Organization

- ▶ 5 people — 2 technical (engineering + security)
- ▶ 160 internal AI agents automating product development
- ▶ GCP-based infra (Brazil region) with on-prem delivery muscle

What Inception unlocks

- ▶ **2026 H2** — on-premise inference GA (TensorRT-LLM); 3 institutional pilots (justice, health, finance).
- ▶ **2027** — multi-tenant sovereign AI serving; marketplace compliance modules; 25+ institutions.

We ask Inception for:

GPU discounts and DGX Cloud access for on-prem appliance certification · TensorRT-LLM/Triton engineering support · DLI training for the technical team · co-marketing to Brazil's regulated institutions.

Let's make regulated AI trustworthy — together.

Beanstech — Beans Tech Inova Simples (I.S.) · São Paulo, Brazil

matheus@beanstech.com.br · beanstech.com.br · linkedin.com/in/mfeijao · (11) 96650-7100